

# Python for Data Science

## Machine Learning 2

**Decision Trees**

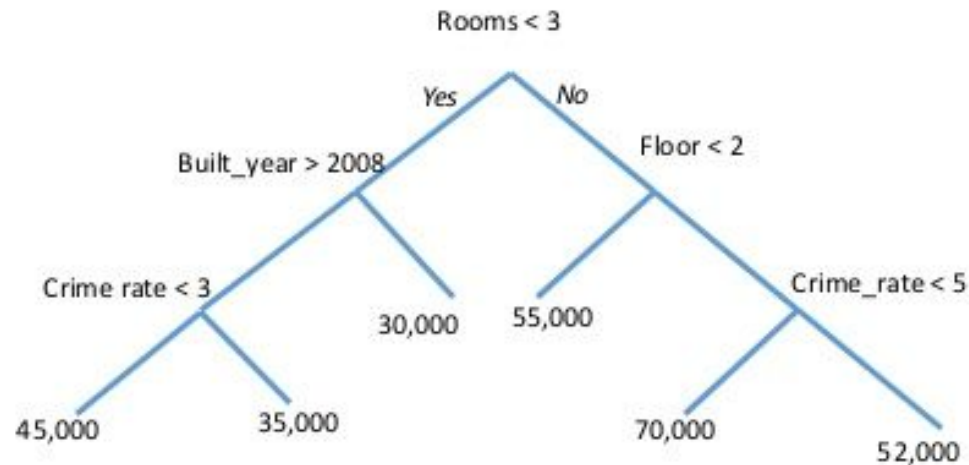
Ensemble Trees

Classification

Wrap Up

# Decision Trees

- ❑ Predicting Model that is based on a tree
- ❑ “Simplified algo”:
  - Select most correlated variable (to target)
  - Select where to split (which value) in order to maximize **separation**
  - Decide when to stop



<http://scikit-learn.org/stable/modules/tree.html>

# Decision Trees

We can split until every observation is separated in the tree

- But should we?
- What will be the RSS?
  - For the training set?
  - For the testing set?

# Decision Trees

We can split until every observation is separated in the tree

- But should we?
- What will be the RSS?
  - For the training set?
  - For the testing set?

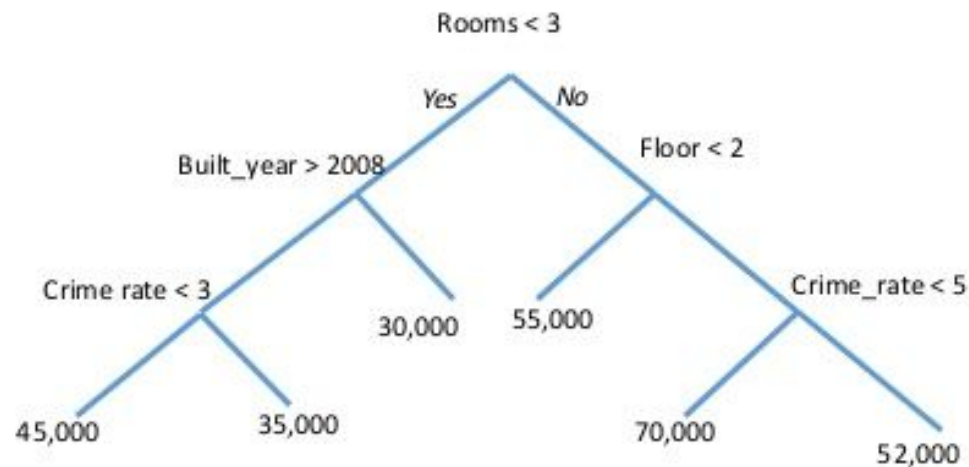


Overfitting!

# Decision Trees

## □ How to use it?

- Decision tree is in fact a list of conditions (if ... Then ... Else ...)
- Follow the steps of the tree to get a prediction



<http://scikit-learn.org/stable/modules/tree.html>

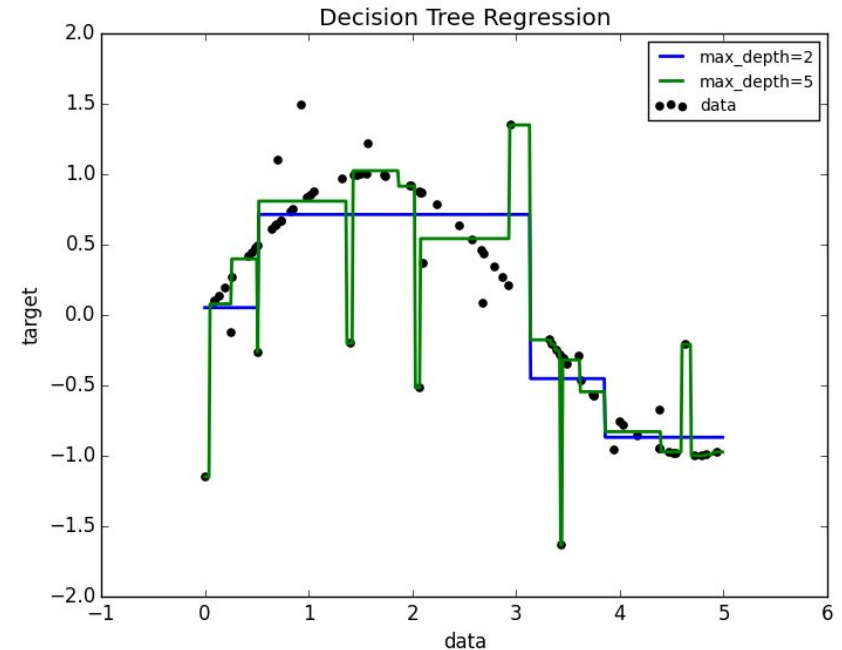
# Decision Trees

## ❑ Pros:

- Easy to understand and interpret
- Testing is easy and quick

## ❑ Cons:

- Bad at generalising (overfitting)
- Unstable
- Training is not obvious:
  - ❑ Approaches: ID3, C4.5, C5.0 and CA



# Decision Tree Regression

```
from sklearn.tree import DecisionTreeRegressor
X = [[0, 0], [2, 2]]
y = [0.5, 2.5]
clf = DecisionTreeRegressor(
    max_depth=10,
    # minimum number of samples required for a split
    min_samples_split=5
)
clf.fit(X, y)
clf.predict([[1, 1]])
```



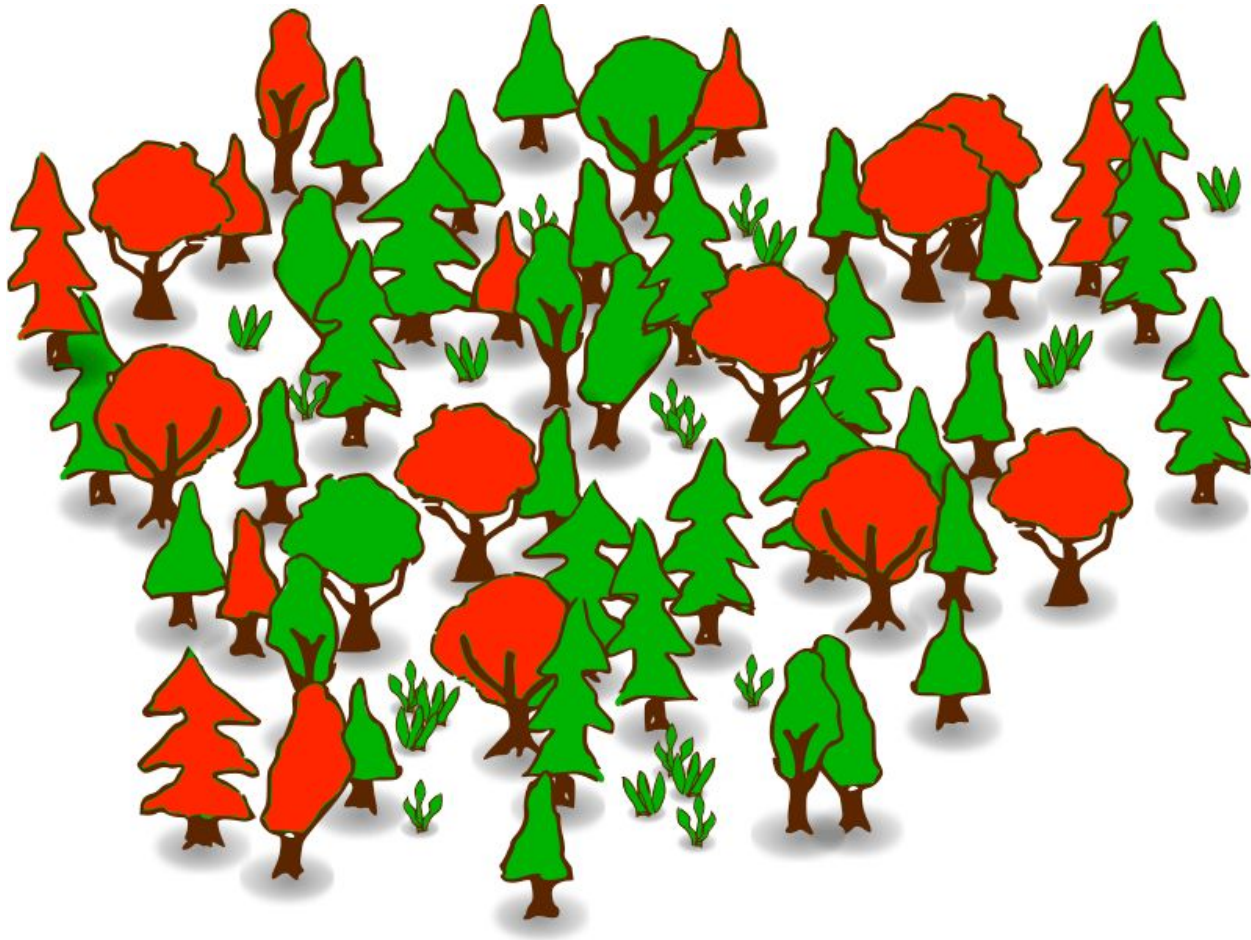
Decision Trees

**Ensemble Trees**

Classification

Wrap Up

# Ensemble Trees



# Ensembling Techniques

## Bagging:

- ❖ **Sample** the input data (features) to generate multiple sets of input data.
  - Usually done with replacement
  - Size of samples is similar to the original data
    - Useful to reduce variance without increasing the bias



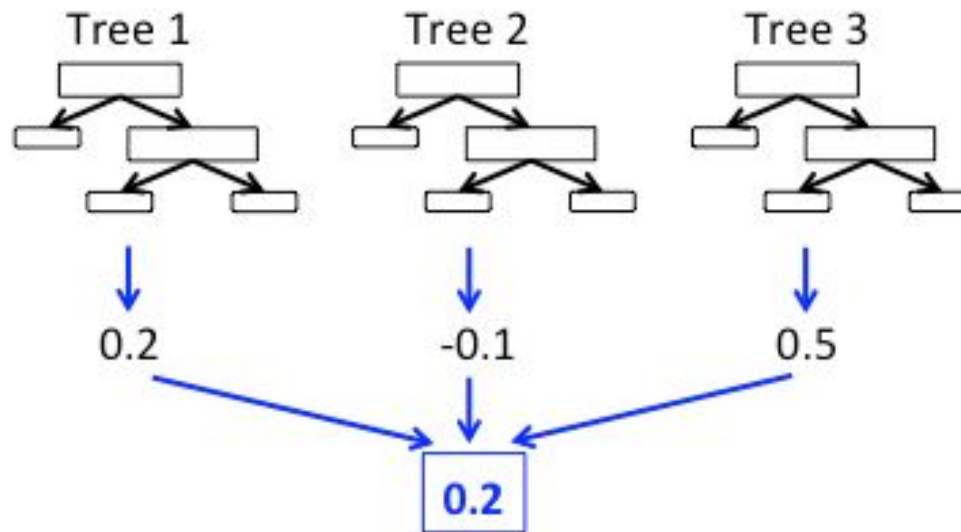
Frodo Baggins

## Boosting

- ❖ Mix a set of **weak** learners to build a strong learner
  - Learners are simple (example decision tree stump)
  - Each weak learner has low variance but high bias
    - Useful to reduce bias without increasing variance

# Random Forest

- ❖ Bagging is similar to an “artificial” increase in the training set:
  - We create  $n$  decision trees, We train each decision tree separately:
    - Training data: draw a sample subset from the training data. Sub-sample size is same as the original input sample size but samples are drawn with replacement
    - Select randomly  $d$  features (usually  $d=\sqrt{m}$ ) without replacement
    - Splits are minimizing gini or entropy
  - Calculate the mean to give the final prediction



# Random Forests in Scikit-Learn

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(n_estimators=10)
```

```
# Train the model using the training set
```

```
model.fit(X_train, y_train)
```

```
# Predict target for the testing set
```

```
y_hat = model.predict(X_test)
```

# Random Forests

## ❑ Pros:

- Don't overfit
- Easy to tune
- Trees grow in parallel - Fast!
- No overfitting Danger
- Ideal to estimate quickly predictability

## ❑ Cons:

- "Black box" models, Not easy to interpret
- Can be complex to deploy (over SQL, Excel, ...)



**Quick Solution (Good most of the time!)**

# Gradient Boosting

- ❑ The algorithm constructs the trees sequentially (slow learning)
- ❑ Each tree is grown using the information from previous tree
- ❑ Typically the tree depth used is smaller than for Random Forests

# Gradient Boosting

- ❑ The algorithm constructs the trees sequentially (slow learning)
- ❑ Each tree is grown using the information from previous tree
- ❑ Typically the tree depth used is smaller than for Random Forests

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set.
2. For  $b = 1, 2, \dots, B$ , repeat:
  - 2.1 Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to the training data  $(X, r)$ .
  - 2.2 Update  $\hat{f}$  by adding in a shrunk version of the new tree:

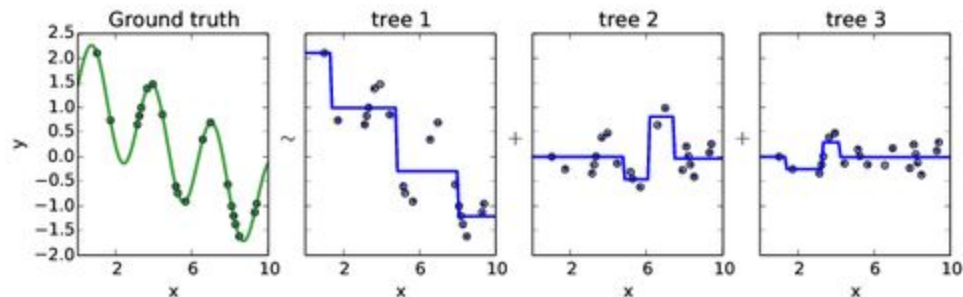
$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x).$$

- 2.3 Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i).$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x).$$





# Gradient Boosting

- Pros:

- Highest performance

- Cons:

- Can't be parallelized (Solution: Parallel implementation such as XGBoost)
- Many parameters to tune (more "Data Scientist time" is necessary)
- Overfitting Danger!
- Same disadvantages as Random Forest



**Highest Performance (Requires more Work!)**

# Gradient Boosting in Scikit-Learn

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
model = GradientBoostingClassifier(n_estimators=100)
```

```
# Train the model using the training set
```

```
model.fit(X_train, y_train)
```

```
# Predict target for the testing set
```

```
y_hat = model.predict(X_test)
```

Decision Trees

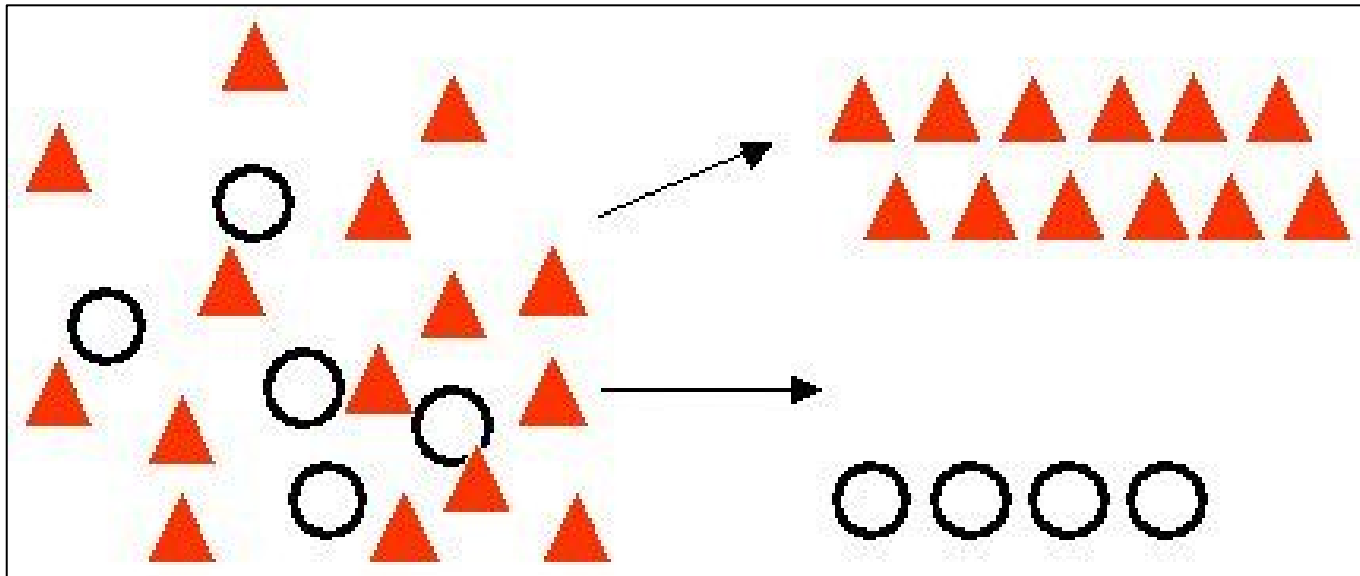
Ensemble Trees

**Classification**

Wrap Up

# Classification

- ❖ Which category a new observation belongs to?



# Evaluation

## ❖ Confusion Matrix - Classification **performance** report

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_true, y_pred)
```

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

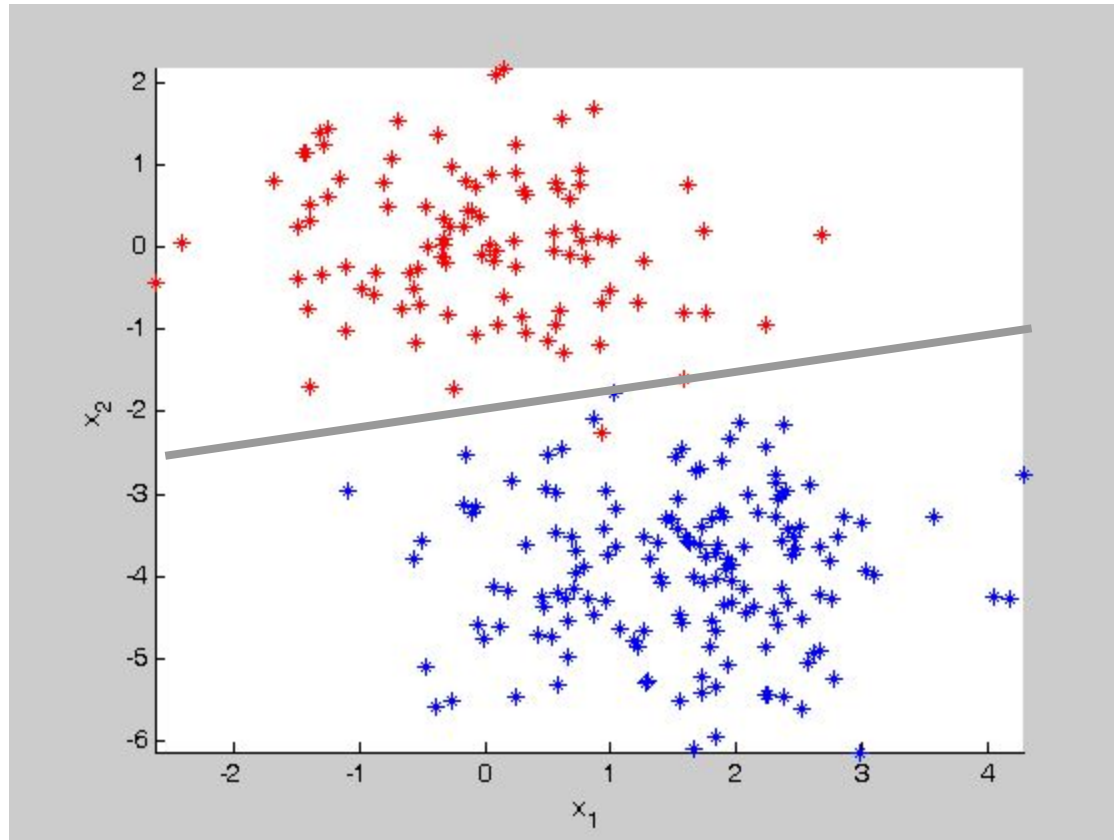
## ❖ Accuracy - How often the prediction is correct?

$$Accuracy = (TP + TN) / Total$$

```
from sklearn.metrics import accuracy_score  
accuracy_score(y_true, y_pred)
```

# Logistic Regression

- ❖ Finding the linear curve that separates the observations into classes:



# Logistic Regression in Scikit-Learn

```
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression(penalty='l2', C=1.0)
```

```
# Train the model using the training set
```

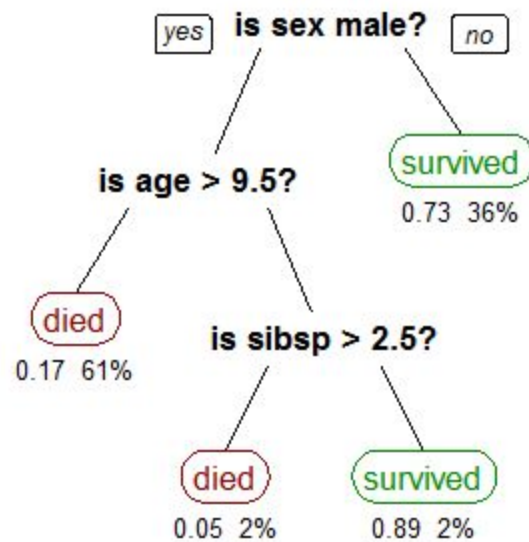
```
logreg.fit(X_train, y_train)
```

```
# Predict target for the testing set
```

```
y_hat = logreg.predict(X_test)
```

# Decision Tree Classifier

- ❖ Finding the decision tree that separates the observations into classes:





# Decision Tree Classifier in Scikit-Learn

```
from sklearn.tree import DecisionTreeClassifier
```

```
model = DecisionTreeClassifier(max_depth=10)
```

```
# Train the model using the training set
```

```
model.fit(X_train, y_train)
```

# Random Forest Classifier in Scikit-Learn

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(  
    n_estimators=100,  
    max_depth=2  
)
```

*# Train the model using the training set*

```
model.fit(X_train, y_train)
```

# Gradient Boosting Classifier in Scikit-Learn

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
model = GradientBoostingClassifier(  
    n_estimators=100,  
    learning_rate=1.0,  
    max_depth=2  
)
```

*# Train the model using the training set*

```
model.fit(X_train, y_train)
```

Decision Tree

Ensemble Trees

Classification

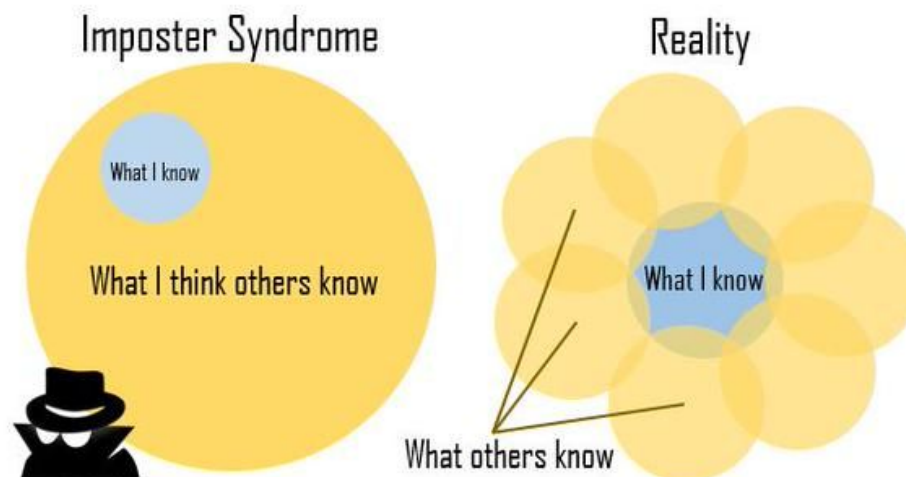
**Wrap Up**

# Wrap-up

- ❑ I hope you think that Python and Machine Learning is **fun** :)
- ❑ To improve your ML skills, you should mix **theory** (MOOCs & Books) with **Kaggle** competitions / projects at **work**.
- ❑ ML is just a **part** of a Data Scientist's job - sometimes short and long-awaited ...
- ❑ You will rarely need to **implement** algorithms (unless you do research) - you must know how to **use** algorithms and **understand** how they work.
- ❑ Remember: Feature engineering is key

# A few more tips...

- ❑ Join a **meetup** or create one
- ❑ Have a **TODO** list for technical books, moocs and articles
- ❑ Use **Twitter** for work
- ❑ Use **wasted time** (e.g. in public transportation) to learn stuff
- ❑ Don't forget that Data Scientists suffer from **imposter syndrome** too...



# Thanks a lot!

[kkarp@equancy.com](mailto:kkarp@equancy.com)

